

Implement K-Nearest Neighbors algorithm on diabetes.csv dataset. Compute confusion matrix, accuracy, error rate, precision and recall on the given dataset.

Dataset link : <https://www.kaggle.com/datasets/abdallamahgoub/diabetes>

```
import pandas as pd
import numpy as np
from sklearn import metrics
```

```
df = pd.read_csv('diabetes.csv')
df
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..	...	...	...	...	...	...
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	Pedigree	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

```
[768 rows x 9 columns]
```

```

df.shape
(768, 9)

# checking for null values
df.isnull().any().value_counts()

False      9
dtype: int64

df.columns

Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
       'Insulin',
       'BMI', 'Pedigree', 'Age', 'Outcome'],
      dtype='object')

df_x = df.drop(columns='Outcome', axis=1)
df_y = df['Outcome']

# When your data has different values, and even different measurement
# units, it can be difficult to compare them.
# The standardization method uses this formula:
#  $z = (x - u) / s$ 

# Where z is the new value, x is the original value, u is the mean and
# s is the standard deviation.

from sklearn.preprocessing import StandardScaler
scale = StandardScaler()
scaledX = scale.fit_transform(df_x)

# split into train and test
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(scaledX, df_y,
                                                    test_size=0.2, random_state=42)

# KNN
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=7)

knn.fit(x_train, y_train)
y_pred = knn.predict(x_test)

# Confusion matrix
cs = metrics.confusion_matrix(y_test, y_pred)
print("Confusion matrix: \n", cs)

Confusion matrix:
[[78 21]
 [28 27]]

```

```

# Accuracy score
ac = metrics.accuracy_score(y_test, y_pred)
print("Accuracy score: ",ac)

# Error rate (error_rate = 1- accuracy)
er = 1-ac
print("Error rate: ",er)

# Precision
p = metrics.precision_score(y_test,y_pred)
print("Precision: ", p)

# Recall
r = metrics.recall_score(y_test,y_pred)
print("Recall: ", r)

# Classification report
cr = metrics.classification_report(y_test,y_pred)
print("Classification report: \n\n", cr)

```

```

Classification report:

```

	precision	recall	f1-score	support
0	0.74	0.79	0.76	99
1	0.56	0.49	0.52	55
accuracy			0.68	154
macro avg	0.65	0.64	0.64	154
weighted avg	0.67	0.68	0.68	154